

search-engine-indexer

search-engine-indexer is a tool for working with search term log files.

Dependencies

search-engine-indexer can be run either from a JAR file or from a Docker image. We recommend running it from a JAR file for performance reasons.

JAR file dependencies

- Java

Docker image dependencies

- Docker

Usage

search-engine-indexer 0.2.0-SNAPSHOT

search-engine-indexer is a tool for working with search term log files

Usage

```
java -jar search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar [SUBCOMMAND] [OPTIONS]
```

Options

```
-v, --version  Show version
-h, --help     Show help
```

Subcommands

generate

```
--dictionary-file  DICTIONARY_FILE
--number-of-output-files  NUMBER_OF_OUTPUT_FILES
--human-bytes-to-write  HUMAN_BYTES_TO_WRITE
--output-directory  OUTPUT_DIRECTORY
```

process

```
--input-directory  INPUT_DIRECTORY
--output-file       OUTPUT_FILE
--maximum-terms-in-memory  MAXIMUM_TERMS_IN_MEMORY
```

JAR file

1. Download

```
wget https://s3.amazonaws.com/search-engine-indexer/jars/search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar
```

2. Run

```
java -jar search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar --help
```

Docker image

Based on `openjdk:11.0.4-jre-slim` which is 204MB.

```
docker run mpereira/search-engine-indexer:0.2.0 --help
```

Implementation details

The `search-engine-indexer` has two main functionalities:

1. generating random unsorted search term log files
2. combining unsorted search term log files into a single alphabetically sorted search term log file

This section will focus on the latter.

The `search-engine-indexer.processor` namespace in `src/search_engine_indexer/processor.clj` has a `process-directory` function that does all the work. It takes as input an *input directory* that should contain search term log files and an *output file* path in which the alphabetically sorted search terms will be created. It also takes a *maximum terms in memory* parameter which sets an upper bound on the memory usage of the program. None of the program's in-memory objects can contain information about more than that amount of search terms.

The algorithm is:

For each input file

For each line (search term) in the input file

Increment in-memory sorted hash map entry for search term

If in-memory sorted hash map has more search terms than maximum

Flush in-memory sorted hash map to search term count files in disk

Aggregate in-disk flushed search term count files

For each aggregated search term count file, in alphabetical order

Append the search term "count" times to the output file

First, it streams each input file sequentially (only one line will be in memory at a time) and accumulates counts in an in-memory sorted hash map. The sorted hash map keys are search terms and values are occurrence counts.

Example

```
{"pinakotek" 2  
 "beer" 1}
```

The program keeps track of the total occurrence counts (in the case above it would be 3) and flushes search term count files to disk when the in-memory sorted hash map is full. Flushing the above sorted hash map would result in an empty in-memory sorted hash map (`{}`) and the following two files on disk:

`pinakotek` contents:

```
pinakotek  
2
```

`beer` contents:

```
beer  
1
```

Assuming the in-memory sorted hash map is updated again:

```
{"pinakotek" 3  
 "ludwig" 5}
```

Flushing it would result in the following disk on disk:

`pinakotek` contents:

```
pinakotek  
2  
3
```

`beer` contents:

```
beer  
1
```

`ludwig` contents:

```
beer  
5
```

- The `pinakotek` file got appended to with a 3
- The `beer` file stayed the same since there were no `beer` occurrences in the sorted hash map
- The `ludwig` file was created with one occurrence count of 5

After all intermediate per-search-term state files are written to disk, their counts are aggregated (summed) and written to disk again. The result of aggregating the files above would be:

pikakotek contents:

```
pinakotek
5
```

beer contents:

```
beer
1
```

ludwig contents:

```
beer
5
```

Disk reads and writes are all sequential and writes are all appends, which should yield high I/O throughput. Multiple intermediate state files are written to, which may cause disk seeks. This can be ameliorated by increasing the *maximum terms in memory* parameter, which will cause fewer intermediate state file writes.

Running a simulation

Even though you're free to clone the repository and build artifacts yourself, in this section we'll make use publicly available AWS S3 objects. The only requirements to follow through the steps are Java and UNIX utilities.

Even though there is a Docker image available for the search-engine-indexer, we'll run the JAR instead so as to not observe reduced IO performance due to doing operations inside a container.

Create a directory for the simulation

```
mkdir search-engine-indexer-simulation
```

1. cd into it

```
cd search-engine-indexer-simulation
```

Download dictionary file

```
wget https://s3.amazonaws.com/search-engine-indexer/dictionaries/german_beer_brands.txt
```

1. Verify dictionary file size

```
wc -l german_beer_brands.txt
```

```
77 german_beer_brands.txt
```
2. Verify dictionary file contents

```

head german_beer_brands.txt

aktienbrauerei kaufbeuren
allgäuer brauerei
asgaard
augustiner-bräu
bayerische staatsbrauerei weihenstephan
berg brauerei
berliner pilsner
berliner weisse
bitburger brauerei
blue girl beer

```

Generate random search term log files

This command will create 10 unsorted search term log files in the `beer_brand_search_terms` directory. Their combined size will be around 500MiB and their contents will come from the `german_beer_brands.txt` dictionary.

```

java -jar search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar \
  generate \
  --dictionary-file german_beer_brands.txt \
  --number-of-output-files 10 \
  --human-bytes-to-write 500MiB \
  --output-directory beer_brand_search_terms

```

Dictionary file:	german_beer_brands.txt
Number of unsorted search term log output files:	10
Human-readable number bytes to write across output files:	500MiB
Output directory for unsorted search term log output files:	beer_brand_search_terms

```

Read dictionary file with 77 terms
Created output directory '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms'
Writing 500MiB across 10 output files
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_0.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_1.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_2.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_3.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_4.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_5.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_6.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_7.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_8.txt'
Created '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search_terms/beer_brand_search_terms_9.txt'
524288119 bytes written in 9.52 seconds (55.06 MB/s)

```

Verify generate output: the beer_brand_search_terms directory

1. Size

```
du -ah beer_brand_search_terms | sort -h
```

```
52M beer_brand_search_terms/0.log
52M beer_brand_search_terms/1.log
52M beer_brand_search_terms/2.log
52M beer_brand_search_terms/3.log
52M beer_brand_search_terms/4.log
52M beer_brand_search_terms/5.log
52M beer_brand_search_terms/6.log
52M beer_brand_search_terms/7.log
52M beer_brand_search_terms/8.log
52M beer_brand_search_terms/9.log
514M  beer_brand_search_terms
```

Looks like the combined size of all generated files is 514M as reported by du. Close enough!

2. beer_brand_search_terms/0.log contents

```
head beer_brand_search_terms/0.log
```

```
bitburger brauerei
hofbräuhaus traunstein
hasseröder
störtebeker braumanufaktur
einbecker brauerei
allgäuer brauerei
brauerei gold oxsen
mecklenburgische brauerei lübz
brauerei gold oxsen
bayerische staatsbrauerei weihenstephan
```

Looks random enough. Let's take a look at another generated file just to make sure it's really random.

3. beer_brand_search_terms/1.log contents

```
head beer_brand_search_terms/1.log
```

```
janssen & bechly brauerei
paulaner brauerei
löwenbräu brauerei
veltins brauerei
list of brewing companies in germany
kronen
janssen & bechly brauerei
```

```
hasseröder
st. erhard brauerei
privatbrauerei wittingen
```

Alright, looks good to me.

Process input directory with randomly generated search term log files (beer_brand_search_terms)

Now that we have a bunch of unsorted search term log files we can combine them into an alphabetically sorted search terms log file.

This command will read the unsorted search term log files in the `beer_brand_search_terms` directory and write their alphabetically sorted search terms to `beer_brand_search_terms.log` while loading at most 100000 (one hundred thousand) search terms (and even "search term"-derived data) in any of the program's in-memory objects.

`maximum-terms-in-memory` sets a hard upper bound on the memory usage of the program. Low values will result in having to write intermediate state to disk more often, but with the advantage of using less memory. High values will result in higher memory usage but less frequent disk writes.

The program will print a dot (.) every time it writes intermediate state to disk.

```
java -jar search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar \
  process \
  --input-directory beer_brand_search_terms \
  --output-file beer_brand_search_terms.log \
  --maximum-terms-in-memory 100000
```

Input directory with unsorted search term log files:	beer_brand_search_terms
Output file to be created with sorted search terms:	beer_brand_search_terms.log
Maximum number of search terms that will be loaded in memory:	100000

```
Processing files in '/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search'
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
/Users/murilo/git/search-engine-indexer/search-engine-indexer-simulation/beer_brand_search
.....
.....
```

```

.....
.....
Term occurrences across all files in '/Users/murilo/git/search-engine-indexer/search-engine
bitburger brauerei 361785
rhanerbräu 361396
brauerei kaiserdom 361197
kaiser bräu 361340
janssen & bechly brauerei 361436
allgäuer brauerei 361344
maisel brau bamberg 361833
warsteiner 361298
flensburger brauerei 361238
grafenwalder 361759
riegele 361593
stadtbrauerei spalt 361926
herrenhäuser brauerei 362955
cölner hofbräu früh 361246
staatliches hofbräuhaus in münchen 362137
köstritzer 361556
vitamalz 362119
diebels 362229
gaffel becker & co 362011
könig brauerei 361475
kulmbacher brauerei 363133
ganter brauerei 362291
koblenzer brauerei 361592
henninger brauerei 361975
radeberger brauerei 362361
fucking hell 361826
eck brauerei 362925
kronen 361306
jeveer brauerei 361465
dortmunder actien brauerei 361335
paulaner brauerei 362367
hacker-pschorr brauerei 362352
heinrich reissdorf 361952
augustiner-bräu 361084
pinkus müller 361765
rothaus 361821
fürstenberg brauerei 362942
hofbräuhaus traunstein 361885
hofbrauhaus arolsen 361563
brauerei gold oxsen 361317
krombacher brauerei 362392
asgaard 361576
oettinger brauerei 361047

```


berliner weisse 361244
privatbrauerei wittingen 362169
list of brewing companies in germany 362604
brauhaus am kreuzberg 362427
st. pauli girl 362812
wernesgrüner 360750
st. erhard brauerei 361367
brauerei zur malzmühle 362470
berliner pilsner 360501
berg brauerei 362180
dortmunder export 360860
veltins brauerei 361039
löwenbräu brauerei 361902
klosterbrauerei andechs 360792
gasthof herold 362595
bayerische staatsbrauerei weihenstephan 362417
könig ludwig schlossbrauerei 362034
bolten-brauerei 361136
störtebeker braumanufaktur 362091
hasseröder 361948
aktienbrauerei kaufbeuren 361275
kuchlbauer brauerei 361407
brauerei gebr. maisel 362642
blue girl beer 361557
freiberger brauhaus 362633
mecklenburgische brauerei lübz 362538
spaten-franziskaner-bräu 361510
herzoglich bayerisches brauhaus tegernsee 360992
licher brauerei 361457
holsten pils 362102
einbecker brauerei 361594
zötler brauerei 361620
erdinger 361485
g. schneider & sohn 361815
Wrote sorted search terms to '/Users/murilo/git/search-engine-indexer/search-engine-indexer-
533332638 bytes (27858110 search terms) processed in 30.02 seconds (17.76 MB/s)

We can see both the per-"search term" and total occurrence counts in the command output. The total number of search terms processed was 27858110.

Verify process output: beer_brand_search_terms.log

1. Size

```
du -ah beer_brand_search_terms.log
```

```
513M    beer_brand_search_terms.log
```

2. Content

```
head beer_brand_search_terms.log
```

```
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren  
aktienbrauerei kaufbeuren
```

This is good evidence that the file is actually alphabetically sorted.

3. Total search terms count

```
cat beer_brand_search_terms.log | wc -l
```

```
27858110
```

27858110 (~27.8 million) is the number of search terms shown in the end of the output for the "Process input directory" step, so this looks good.

4. Unique search terms count

```
cat beer_brand_search_terms.log | uniq | wc -l
```

```
77
```

77 is the number of terms in the dictionary, so this is also looking good.

5. Unique search term counts

```
uniq -c beer_brand_search_terms.log | head
```

```
361275 aktienbrauerei kaufbeuren  
361344 allgäuer brauerei  
361576 asgaard  
361084 augustiner-bräu  
362417 bayerische staatsbrauerei weihenstephan  
362180 berg brauerei  
360501 berliner pilsner  
361244 berliner weisse  
361785 bitburger brauerei  
361557 blue girl beer
```

It seems that the search terms were written in order, otherwise `uniq -c` wouldn't have worked. Also, comparing the output above with the output

produced in the "Process input directory" step should demonstrate that search terms were written correctly.

That's it

In this simulation we generated random unsorted search term log files, combined them into an alphabetically sorted search term log file and verified that the outputs were correct. Feel free to run simulations with different parameters!

Just for fun, let's try a much smaller `maximum-terms-in-memory` value. Let's set it to 1000. We'd expect reduced throughput given more frequent intermediate state writes to disk.

```
java -jar search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar \
  process \
  --input-directory beer_brand_search_terms \
  --output-file beer_brand_search_terms.log \
  --maximum-terms-in-memory 1000 \
  | grep 'processed in'
```

533332638 bytes (27858110 search terms) processed in 233.20 seconds (2.29 MB/s)

We can see one order of magnitude less throughput (~17 MB/s -> ~2 MB/s) by decreasing `maximum-terms-in-memory` three orders of magnitude (100000 -> 1000).

Let's see what happens when setting it to a much higher value, 10000000 (10 million). Based on the total number of search terms we've seen above (~27.8 million) the in-memory buffer will be flushed to disk at most thrice.

```
java -jar search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar \
  process \
  --input-directory beer_brand_search_terms \
  --output-file beer_brand_search_terms.log \
  --maximum-terms-in-memory 10000000 \
  | grep 'processed in'
```

533332638 bytes (27858110 search terms) processed in 28.49 seconds (18.72 MB/s)

Only a marginal throughput improvement from 100000. Likely means that at those parameter sizes the bottleneck isn't in intermediate state file disk writes.

Development

Work on `search-engine-indexer` is mostly done on Emacs. The workflow looks like:

1. A CIDER session is started with `M-x cider-jack-in`

2. Code is evaluated with `cider-eval-sexp-at-point` or `cider-eval-buffer`
3. Tests are run with `cider-test-run-test` or `cider-test-run-ns-tests`

Dependencies

- Java
- Leiningen
- Docker

Check out repository

```
git clone git@github.com:mpereira/search-engine-indexer.git
```

cd into repository

```
cd search-engine-indexer
```

Running tests

```
lein test
```

Building uberjar

```
lein do clean, uberjar
```

Publishing uberjar

```
aws s3 cp --acl public-read \  
  target/uberjar/search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar \  
  s3://search-engine-indexer/jars/search-engine-indexer-0.2.0-SNAPSHOT-standalone.jar
```

Building Docker image

```
docker build -t mpereira/search-engine-indexer:0.2.0 .
```

Publishing Docker image

```
docker login
```

```
docker push mpereira/search-engine-indexer:0.2.0
```

License

Copyright © 2019 Murilo Pereira <murilo@murilopereira.com>

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.